Certified Should Mean Checked

Keeping Genie Agent certification honest after the agent changes

Shatakshi Chaudhari
MS in Software Management, Carnegie Mellon University
Previously Product Management Intern, Cisco
shataksc@andrew.cmu.edu
LinkedIn · Product portfolio

Key
T1–T10: tests I ran, with full results in the Test Log
R1–R9: outside sources, also listed in the Test Log
H1–H3: assumptions · **P0–P2:** requirement priority

TL;DR

Analysts certify Genie Agents so business users know which ones to trust. But the badge has nothing behind it, and it doesn't notice when the agent changes. In my trial I certified an agent, then told it to round revenue to the nearest $1,000. The benchmarks failed. The badge stayed Certified. Then Improve with Genie swapped out two of my benchmarks, and it was still Certified.

The fix: **Certified means "checked, on this version."** Save who certified which version, and what the results were. Turn the badge amber when someone makes an edit that can change answers. Then give the owner a one-click re-check that compares result tables with the certified version, with no LLM doing the grading. Now is the time: Databricks wants everyone building and sharing agents, AI helpers now make many of the edits, and Genie stops being free on January 31, 2027. Success is 80% of certified agents still on their certified version after 90 days.

> **Prototype:** https://shatakshi107.github.io/genie-certification-prototype/

Problem, Opportunity, and Evidence

Who and when: Priya (fictional) is a BI analyst at Harbor Health. She builds a Genie Agent on governed tables, certifies it, and shares it with a sales team. Afterward she, a teammate, or Genie's AI helper keeps editing it, while business users keep picking it because it says Certified. Priya wants it to stay trustworthy as it improves. The sales team wants to act on an answer without checking it themselves.

	One editor	Many editors
High reliance	**Solo analyst · Second** A team relies on it, but one person maintains it.	**Priya · Priority** Enterprise workspace. A shared, certified agent, edited by several people and the AI helper.
Low reliance	**Personal sandbox · Out of scope** Little to protect.	**Team experiment · Future** Edited often, but few rely on it yet.

Priority segment for the beta.

> ### The benchmark caught my change. The badge didn't.
>
> Databricks says a certified asset "has met internal standards for accuracy, completeness, and trust" (Databricks docs). In practice it's a manual tag with no version attached, so it survives any edit. Keep the promise, and tie it to the version someone actually checked.

Friction	Observed in my trial	Test
No evidence behind the badge	Assign certification is a choice of Certified, Deprecated or None, with **no score, date or version.** The tooltip says only "This asset is certified."	T8, T1
The badge never lapses	As a deliberate test, I added "Always round revenue to the nearest $1,000." The Agents list showed **Certified beside "Modified 22 seconds ago."**	T1, T2
Scores aren't reliable enough	On 8 benchmark questions, Agent and Chat mode each scored **4/8, with different questions passing.** Verdicts also flipped between runs. Correct answers failed for "Extra Rows." The sample Starter Space scored **0/4 with every answer correct.**	T2, T4
AI changes without approval	I skipped Improve with Genie's question. It **replaced 2 benchmarks on the certified agent,** then reported "5 out of 8 passing, up from 4" across different question sets. In other runs, approval prompts expired in about 30–40 seconds, and 1 of 3 still changed the agent.	T2, T5
No trail or way back	The audit log records the **AI's edits under my email,** with IDs and no text. There is **no version history or revert.**	T7, T8

This grading noise is a known LLM-judge problem, not a Genie bug. LLM judges disagree with themselves across repeated runs (R1), and a production text-to-SQL judge reached only $\kappa = 0.42$ agreement with human labels on a random sample (R2). My tests prove it happens. The beta measures how often.

ID	Assumption to validate	Beta test
H1	Business users read Certified as "correct today."	A short survey
H2	Owners re-check within a week when a badge lapses.	The re-check metric below
H3	Certified agents are edited often enough to matter.	Count post-certification edits in the audit log

Why now

Scale	PepsiCo is moving about 30k reports to about 50 Genie agents, and GM runs hundreds (Data + AI Summit 2026 keynote), and Genie Code writes about 60% of LakeFlow pipelines. More people and AI will edit trusted agents.
AI reach	Since June 2026, Genie One prioritizes certified assets (Databricks release notes), so a stale badge also steers AI answers.
Timing	Genie usage is free through January 31, 2027 (Databricks release notes), so trust built now will affect usage once billing starts.

Why it's worth doing: Fewer wrong numbers reach leadership. Genie is billed by use ($0.07/DBU), so trust turns directly into queries. Power BI and Snowflake both let teams endorse or verify content, but from their docs, neither shows what changed in the answers (R9). It also reuses benchmarks people already wrote.

Jobs to Be Done

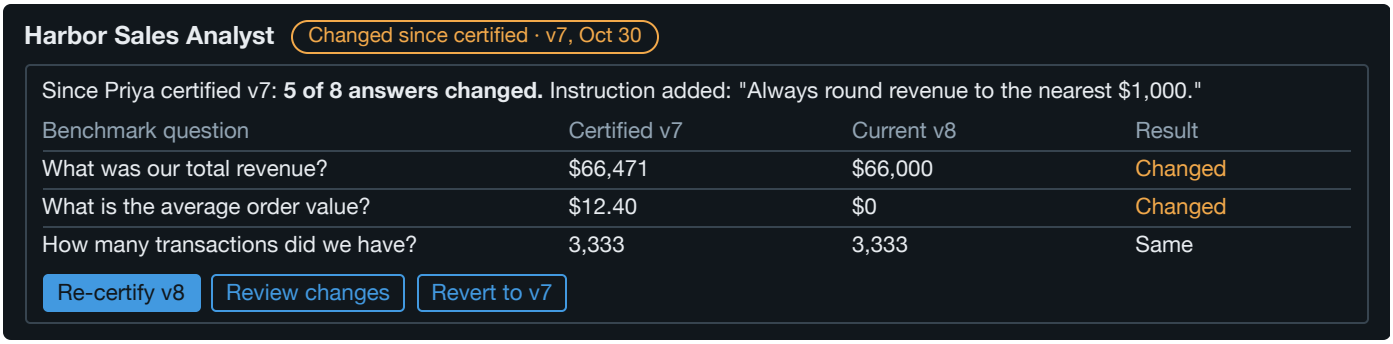
Core job (analyst): ensure an AI data assistant gives my business team correct answers, before and after it is shared.
Desired outcomes, ranked:

1. Minimize the likelihood that business users rely on an agent that changed after it was checked. (P0-2)
2. Minimize the time it takes to see whether a change altered answers people already trust. (P0-3)
3. Minimize the likelihood that tests or settings change without my explicit approval. (P1-2)
4. Minimize the effort to show stakeholders who checked an agent, when, and on what. (P0-1)

Related job (business user): decide quickly whether I can act on an answer without re-checking it.

Proposed Product Experience

Don't ask a noisy grader whether answers are right. Show the analyst what changed and let them decide. You can click through it in the prototype, which has a guided tour.



Sketch of the proposed amber badge and re-check panel. Illustrative numbers.

User journey:

Certify Priya certifies v7. A receipt saves who, when and the results. → **Trust** Users see "Certified · v7 · Oct 30" in the list, Discover and chat. → **Change** A person or AI helper edits it. The badge turns amber. → **Re-check** One click re-runs certified questions and shows what changed. → **Decide** Re-certify v8, keep editing, or preview v7 and revert (P1).

Key decisions

Decision	Why	Trade-off
No score gate	Verdicts flip, so good agents would fail at random.	Relies on human judgment, so I track whether analysts review changes.
Compare tables instead of grading	Comparing two tables gives the same result every time.	It can't say whether a change is good. That call stays with the analyst.
Lapse by type of edit	Instructions, SQL, tables, joins and benchmarks lapse it. Titles and sharing don't.	Fewer false alarms, but someone has to maintain the list.

Amber instead of removing the badge	Removing it would alarm users and hide history.	Users can learn to ignore amber, so a guardrail tracks it.
Benchmarks in v1, live traffic later	Cheaper, and avoids privacy issues.	Only covers those questions.

Top alternative: require a passing benchmark score to certify

	Show what changed (proposed)	Score gate
Strengths	Isn't affected by grader noise. Shows the actual impact. Small, and reuses benchmarks.	Simple, automatic, and familiar from CI/CD.
Weaknesses	Needs a human decision each time. Covers only benchmark questions. Run-to-run noise can look like a change, so P1-4 tags unstable questions.	Blocks good agents when verdicts flip. Owners stop certifying. Invites fixing the test, which Improve with Genie proposed in all 3 runs (T5).

Requirements

ID	Requirement and acceptance test
P0-1	**Receipt:** Certifying stores certifier, time, version, benchmark results and a note. With no benchmarks, the badge says "no tests attached." Test: all appear in the dialog and on hover.
P0-2	**Lapse:** Any listed edit, including by the AI helper, turns the badge amber in the header, Agents list, Discover and chat. Test: editing an instruction turns it amber by the next page load; editing the title doesn't.
P0-3	**Re-check:** One click re-runs certified questions, marks each result table Changed or Same, and shows what changed in the configuration and who changed it. Test: 8 questions finish within 2 minutes, with progress shown, at no more than the cost of one benchmark run.
P1-1	**Revert with preview:** Show certified and current answers side by side before confirming. V1 only reverts to the certified version.
P1-2	**AI approval:** A timed-out or skipped prompt applies no change. AI edits to benchmarks need explicit approval.
P1-3	**Safer "Update ground truth":** Show what it replaces, and keep the original expected answer (T6).
P1-4	**Flag flips, and ask instead of guessing:** Tag questions whose verdict changed across runs or modes. If one is ambiguous, ask the analyst to choose a meaning and save it (R3).
P2	Full version history. Audit entries showing "proposed by Genie Code, approved by [user]." Real-question replay.

Out of scope: grader accuracy, model determinism, draft vs published agents, and object types beyond Genie Agents.

Success Metrics

Type	Metric	Good	Bad	Why this threshold
North star	**Certified-and-current rate: % of certified agents on their certified version (90 days)**	≥ 80%	< 50%	Under half means the badge is usually out of date. The baseline is unknown today.
Launch	% of amber agents re-checked within 7 days	≥ 50%	< 20%	Tests H2. Low means re-check is hard to find or not valued.
Launch	Median time from change to re-check	< 1 day	> 7 days	People keep getting wrong answers until someone checks.
Quality	% of re-certifications with changes reviewed	≥ 60%	< 20%	Low means people are rubber-stamping.
Guardrail	Drop in certified agents (amber counts)	≤ 10%	> 25%	Catches certification becoming a burden.
Guardrail	Lapses per agent per week	< 1	≫ 1	Too many teach people to ignore amber.
Stop rule	Stop if any behavior-changing edit leaves a badge blue. Pause expansion if owners remove certification from more than 25% of lapsed agents within 30 days, compared with similar workspaces outside the preview.			

Measurement and rollout: Certification events and versions come from system tables. Launch as a private preview with 5–10 design partners to set the thresholds, then turn it on for new certifications, then existing ones. Long term, certified agents should get fewer thumbs-down than uncertified ones. **What would change the plan:** if fewer than 20% of amber agents get re-checked within a week, owners don't act on the badge, and the signal moves to a direct alert to the owner instead.